Benchmark Studio

LLM MODEL PERFORMANCE COMPARISON ANALYSIS

Report Milestone 3

Report by:

Akshara Sri Lakshmipathy - akla8196@colorado.edu

Harish Nandhan Shanmugam - hash1366@colorado.edu

Shivaraj Senthil Rajan - shse1502@colorado.edu

In this we imported a CSV file (battles_data_cleaned.csv) into a DataFrame and then displayed the first few rows.

	turn	language	is_code	is_refusal	model	sum_user_tokens	sum_assistant_tokens	context_tokens	dedup_tag_high_freq	dedup_tag_sampled	tstamp_period	information_fulfillment	Math	spec
0	1	French	False	False	gemma-2-2b-it	19	328	19	False	True	2024-08-04	True	False	
1	3	English	False	False	athene-70b-0725	51	1846	1434	False	True	2024-08-04	True	False	
2	6	Italian	False	False	athene-70b-0725	472	3504	3396	False	True	2024-08-04	False	False	
3	2	English	False	False	llama-3-70b-instruct	84	1047	580	False	True	2024-08-04	True	False	
4	1	French	False	False	gpt-4-turbo-2024-04-09	29	556	29	False	True	2024-08-04	False	False	

HANDLING CLASS IMBALANCE

It calculates the frequency of models in the dataset, filter them and keep only with at least 1000 occurrences and create a new filtered DataFrame containing only these models.

```
model
chatgpt-4o-latest                    9099
gpt-4o-2024-08-06                    5395
llama-3.1-405b-instruct              3557
gemini-1.5-pro-exp-0801              3427
gpt-4o-2024-05-13                    3340
llama-3.1-70b-instruct               3321
claude-3-5-sonnet-20240620           3219
mistral-large-2407                   3113
gemini-1.5-pro-api-0514              3030
gpt-4o-mini-2024-07-18               2535
llama-3.1-8b-instruct                2486
reka-core-20240722                   2465
reka-flash-20240722                  2334
athene-70b-0725                      2264
gpt-4-turbo-2024-04-09               2047
claude-3-opus-20240229               1939
gemini-1.5-flash-api-0514            1927
gemma-2-27b-it                       1844
deepseek-v2-api-0628                 1777
gemma-2-2b-it                        1600
llama-3-70b-instruct                 1089
gpt-4-0125-preview                   1002
gpt-4-1106-preview                    991
deepseek-coder-v2-0724                987
...
phi-3-mini-4k-instruct-june-2024      628
gemini-advanced-0514                  219
mixtral-8x7b-instruct-v0.1            188
Name: count, dtype: int64
```

It calculates the value counts of the 'model' column in the filtered data and prints the number of models.

```
22
```

To balance the dataset we can control the number of samples for each unique model. It filters models that have at least 1000 samples, and then resamples each model to ensure an equal representation of samples either by undersampling or oversampling. The code first saves the balanced dataset (balanced_data) to a CSV file and printed the first few rows of the data are displayed to confirm the dataset's successful import in resampled_data.head()

	turn	language	is_code	is_refusal	model	sum_user_tokens	sum_assistant_tokens	context_tokens	dedup_tag_high_freq	dedup_tag_sampled	tstamp_period	information_fulfillment	Math	spec
0	1	Persian	False	False	chatgpt-4o-latest	92	227	92	False	True	2024-08-09	False	False	
1	2	English	False	True	chatgpt-4o-latest	1880	2504	2877	False	True	2024-08-08	True	False	
2	1	English	False	False	chatgpt-4o-latest	8	806	8	False	True	2024-08-10	False	True	
3	1	English	False	False	chatgpt-4o-latest	7	78	7	False	True	2024-08-08	False	False	
4	1	English	False	False	chatgpt-4o-latest	36	83	36	False	True	2024-08-11	True	False	

After this the code initializes and applies a LabelEncoder to the 'language' column from categorical text to numeric values. A LabelBinarizer is used to binarize the target labels, which helps transform the target variable into a format compatible with the multiclass classification. And then code splits the dataset into training and testing sets (80% training, 20% testing).

MODEL IMPLEMENTATION

1. RANDOM FOREST

A **Random Forest** binary classifier is trained, resulting in multiple classifiers. After training, the classifiers are used to collect the probability scores of each instance in the test set for each class and the class with the highest probability is then selected for each test instance where it calculates and displays metrics—accuracy, precision, recall, and F1-score—for each of the binary classifiers on the test set

```
Metrics for class athene-70b-0725:
  Accuracy: 0.982840909090909
  Precision: 0.888235294117647
  Recall: 0.727710843373494
  F1-score: 0.8
Metrics for class chatgpt-4o-latest:
  Accuracy: 0.9526136363636364
  Precision: 0.20270270270270271
  Recall: 0.0189873417721519
  F1-score: 0.034722222222222224
Metrics for class claude-3-5-sonnet-20240620:
  Accuracy: 0.98125
  Precision: 0.8844827586206897
  Recall: 0.6610824742268041
  F1-score: 0.7566371681415929
Metrics for class claude-3-opus-20240229:
  Accuracy: 0.9872159090909091
  Precision: 0.9104258443465492
  Recall: 0.7908163265306123
  F1-score: 0.8464163822525598
Metrics for class deepseek-v2-api-0628:
  Accuracy: 0.9882386363636364
  Precision: 0.9019886363636364
  Recall: 0.8214747736093143
  F1-score: 0.8598510494245092
...
  Accuracy: 0.9841477272727273
  Precision: 0.9288079470198676
  Recall: 0.7038895859473023
  F1-score: 0.8008565310492506
```

Printed the overall accuracy of the Random Forest classifier across all classes

```
Overall Accuracy of Random Forest Classifier: 0.69
```

gpt-4-0125-preview and llama-3-70b-instruct showing high precision, recall, and F1-scores, while models like chatgpt-4o-latest underperformed. Most of the models have moderate performance, with scores ranging from 0.60 to 0.80.

```
Classification Report:
                            precision    recall  f1-score   support

          athene-70b-0725       0.69      0.74      0.71       830
        chatgpt-4o-latest       0.15      0.08      0.11       790
claude-3-5-sonnet-20240620       0.68      0.68      0.68       776
    claude-3-opus-20240229       0.69      0.81      0.75       784
      deepseek-v2-api-0628       0.75      0.84      0.79       773
 gemini-1.5-flash-api-0514       0.76      0.82      0.79       821
   gemini-1.5-pro-api-0514       0.68      0.68      0.68       847
   gemini-1.5-pro-exp-0801       0.67      0.65      0.66       796
            gemma-2-27b-it       0.77      0.82      0.79       822
             gemma-2-2b-it       0.79      0.85      0.82       828
       gpt-4-0125-preview       0.83      0.93      0.88       747
    gpt-4-turbo-2024-04-09       0.72      0.79      0.75       799
         gpt-4o-2024-05-13       0.66      0.65      0.66       785
         gpt-4o-2024-08-06       0.13      0.07      0.09       769
    gpt-4o-mini-2024-07-18       0.72      0.70      0.71       840
      llama-3-70b-instruct       0.83      0.94      0.88       825
   llama-3.1-405b-instruct       0.66      0.61      0.63       807
    llama-3.1-70b-instruct       0.66      0.67      0.67       796
     llama-3.1-8b-instruct       0.72      0.73      0.72       797
        mistral-large-2407       0.66      0.65      0.65       789
        reka-core-20240722       0.71      0.75      0.73       782
       reka-flash-20240722       0.72      0.74      0.73       797
...
 [ 15  15   8  10  10   6   8   5   9   6   1   9   8  21   3   6  12   9
   11  16 585   9]
 [  8   9  14  13   6  11   9   6   8  10   7   9  12  13  11   8  10  14
    9  13  11 586]]
```

The matrix provides a view of how well each model is performing to the others in predicting different classes.

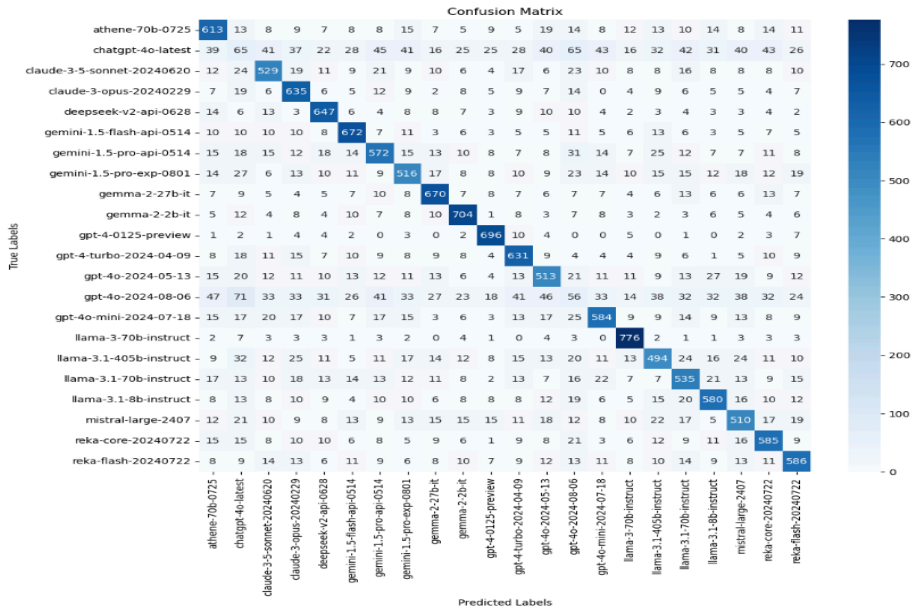

The area under each curve (AUC) is an indicator of how well the model can distinguish between classes, with values closer to 1 indicating better performance.

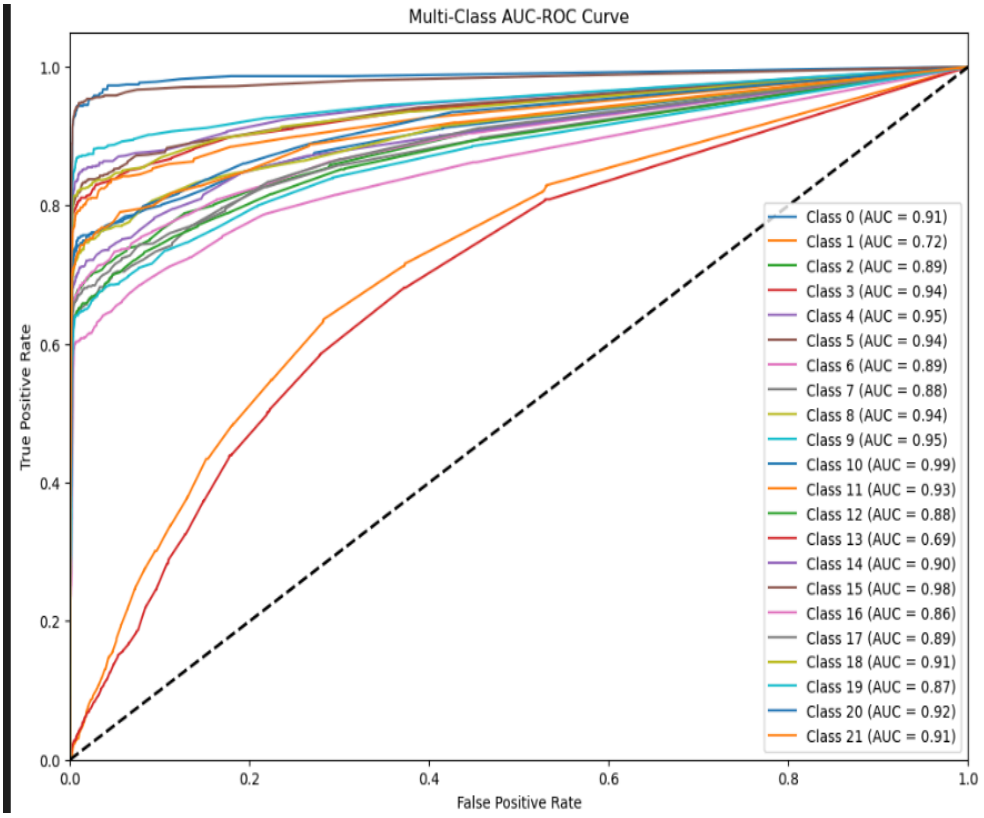

The Random Forest Classifier undergoes hyperparameter tuning using GridSearchCV to find the optimal parameters, which are used to train a One-vs-Rest classification model. After training, predictions are made by selecting the class with the highest probability for each instance. The evaluation involves calculating overall accuracy, classification report for precision, recall, and F1-score, and visualizing the confusion matrix. Additionally, AUC-ROC curves are plotted.

The Random Forest model achieved an overall accuracy of 0.69, with gpt-4-0125-preview and llama-3-70b-instruct performing F1-scores above 0.88. Models like chatgpt-4o-latest had low F1-scores around 0.11.

```
Fitting 3 folds for each of 216 candidates, totalling 648 fits
Best Parameters: {'bootstrap': True, 'max_depth': None, 'min_samples_leaf': 1, 'min_samples_split': 2, 'n_estimators': 300}
Overall Accuracy of Tuned Random Forest Classifier: 0.69
Classification Report:
                           precision    recall  f1-score   support

          athene-70b-0725       0.70      0.74      0.72       830
        chatgpt-4o-latest       0.16      0.09      0.11       790
claude-3-5-sonnet-20240620       0.67      0.68      0.68       776
    claude-3-opus-20240229       0.70      0.81      0.75       784
      deepseek-v2-api-0628       0.75      0.84      0.79       773
 gemini-1.5-flash-api-0514       0.77      0.82      0.79       821
   gemini-1.5-pro-api-0514       0.69      0.68      0.68       847
   gemini-1.5-pro-exp-0801       0.67      0.65      0.66       796
          gemma-2-27b-it       0.77      0.82      0.79       822
           gemma-2-2b-it       0.79      0.85      0.82       828
       gpt-4-0125-preview       0.83      0.93      0.88       747
    gpt-4-turbo-2024-04-09       0.72      0.78      0.75       799
        gpt-4o-2024-05-13       0.67      0.65      0.66       785
        gpt-4o-2024-08-06       0.14      0.08      0.10       769
    gpt-4o-mini-2024-07-18       0.72      0.70      0.71       840
      llama-3-70b-instruct       0.83      0.94      0.88       825
   llama-3.1-405b-instruct       0.65      0.61      0.63       807
    llama-3.1-70b-instruct       0.65      0.67      0.66       796
     llama-3.1-8b-instruct       0.71      0.73      0.72       797
...
                  accuracy                           0.69     17600
                 macro avg       0.67      0.69      0.68     17600
              weighted avg       0.67      0.69      0.68     17600
```

High values on the diagonal elements indicate a high number of correct predictions, while off-diagonal elements show the misclassifications. For example, model 'deepseek-v2-api-0628' has 649 correct predictions for one of its classes, while it has misclassified other classes a few times as shown by the non-diagonal elements.

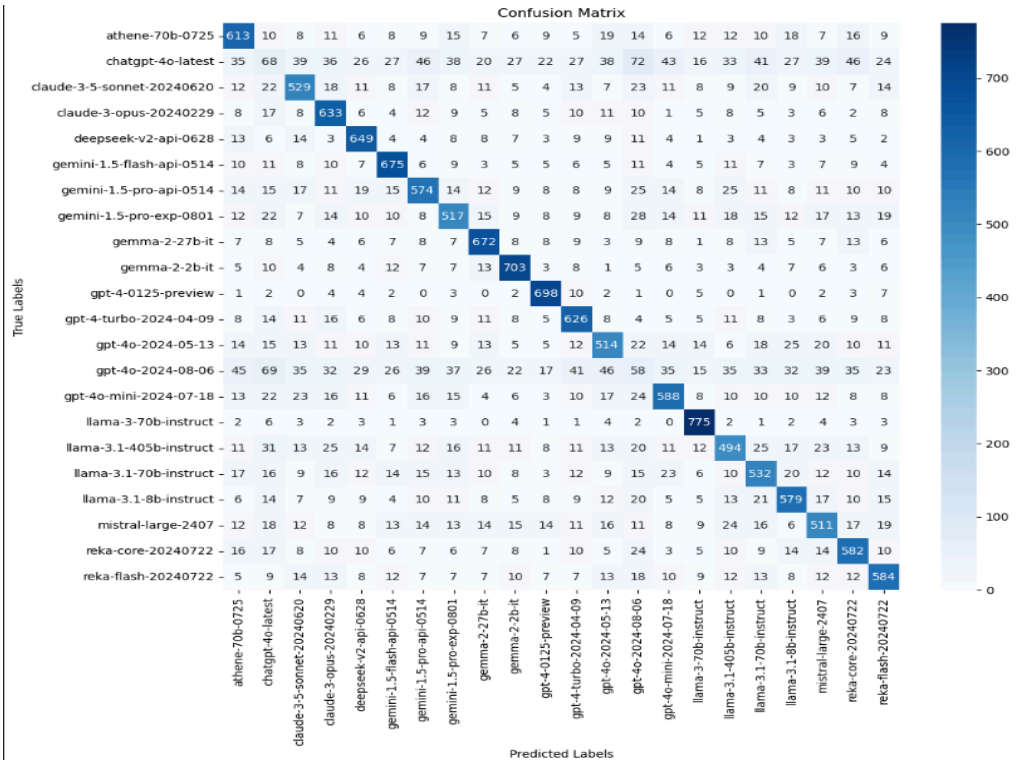

Higher AUC values, closer to 1, indicate better model performance.

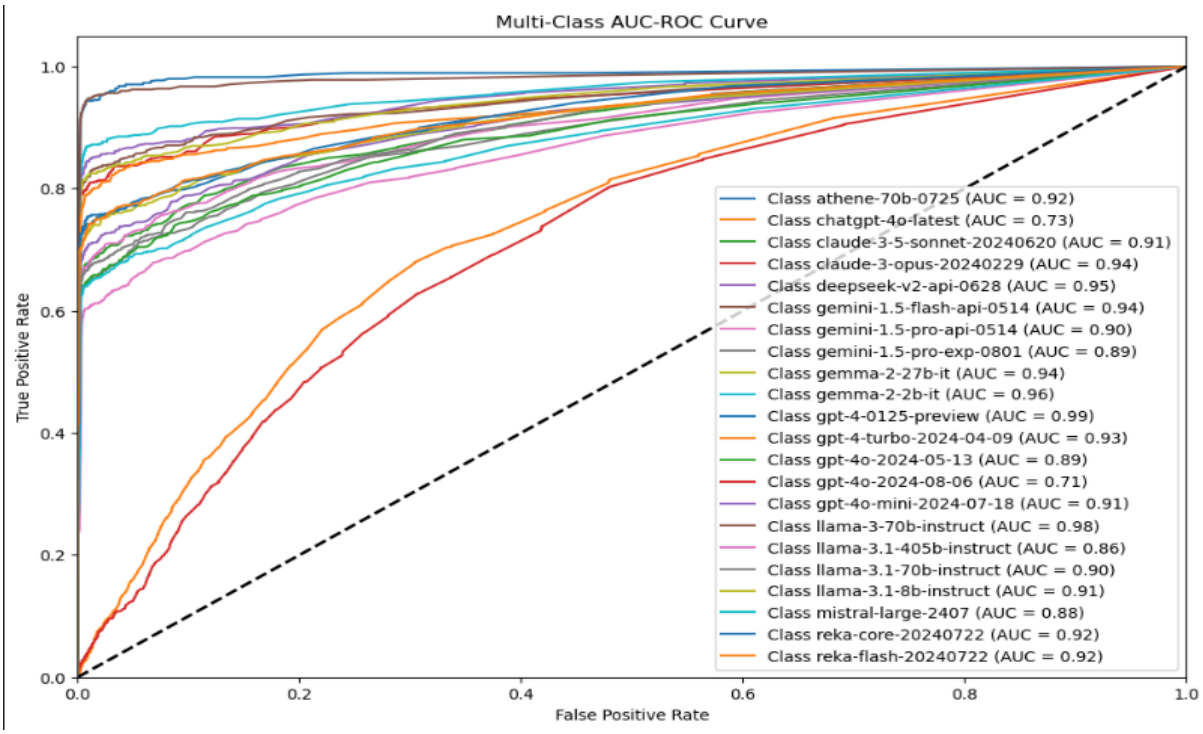

2. XG BOOST

The **XGBoost** classifier achieved an overall accuracy of 0.68.

We have used XGBoost for a multi-class classification using a One-vs-Rest approach and involves the following steps: Data Loading and Preprocessing, Hyperparameter Tuning using RandomizedSearchCV to tune hyperparameters and selecting the best parameters. Then it trains separate One-vs-Rest XGBoost classifiers for each class. Making predictions which involves calculating overall accuracy, classification report for precision, recall, and F1-score, and visualizing the confusion matrix. Additionally, AUC-ROC curves are plotted.

```
Fitting 3 folds for each of 100 candidates, totalling 300 fits
Best Parameters for XGBoost: {'subsample': 1.0, 'scale_pos_weight': 2, 'reg_lambda': 10, 'reg_alpha': 0.1, 'n_estimators': 700, 'min_child_weight': 1, 'max_depth': 15, 'learning_rate':
Overall Accuracy of Tuned XGBoost Classifier: 0.68
Classification Report:
                          precision    recall  f1-score   support

          athene-70b-0725       0.69      0.73      0.71       830
        chatgpt-4o-latest       0.17      0.10      0.12       790
claude-3-5-sonnet-20240620       0.66      0.66      0.66       776
    claude-3-opus-20240229       0.71      0.80      0.75       784
      deepseek-v2-api-0628       0.75      0.80      0.78       773
 gemini-1.5-flash-api-0514       0.74      0.80      0.77       821
   gemini-1.5-pro-api-0514       0.69      0.67      0.68       847
   gemini-1.5-pro-exp-0801       0.64      0.63      0.64       796
           gemma-2-27b-it       0.75      0.81      0.78       822
            gemma-2-2b-it       0.78      0.85      0.81       828
       gpt-4-0125-preview       0.82      0.93      0.87       747
    gpt-4-turbo-2024-04-09       0.71      0.77      0.74       799
         gpt-4o-2024-05-13       0.64      0.63      0.63       785
         gpt-4o-2024-08-06       0.13      0.09      0.10       769
    gpt-4o-mini-2024-07-18       0.71      0.70      0.71       840
      llama-3-70b-instruct       0.82      0.93      0.87       825
   llama-3.1-405b-instruct       0.62      0.59      0.61       807
    llama-3.1-70b-instruct       0.67      0.63      0.65       796
     llama-3.1-8b-instruct       0.70      0.70      0.70       797
...
                 accuracy                           0.68     17600
                macro avg       0.66      0.68      0.66     17600
             weighted avg       0.66      0.68      0.67     17600
```

The diagonal values, which show the highest numbers (e.g., 603 for the first class), indicate correct predictions.

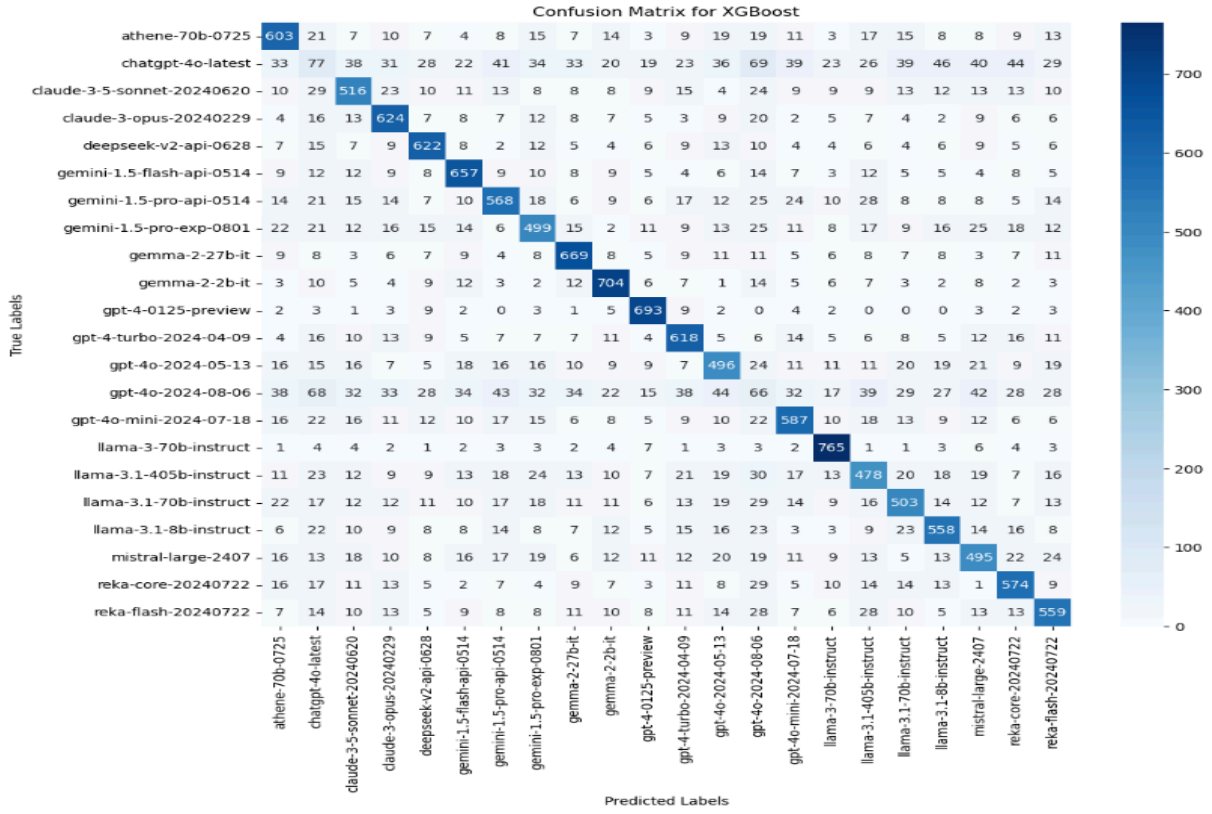

The curves closer to the top-left corner of the plot indicate better performance, with AUC values closer to 1.0

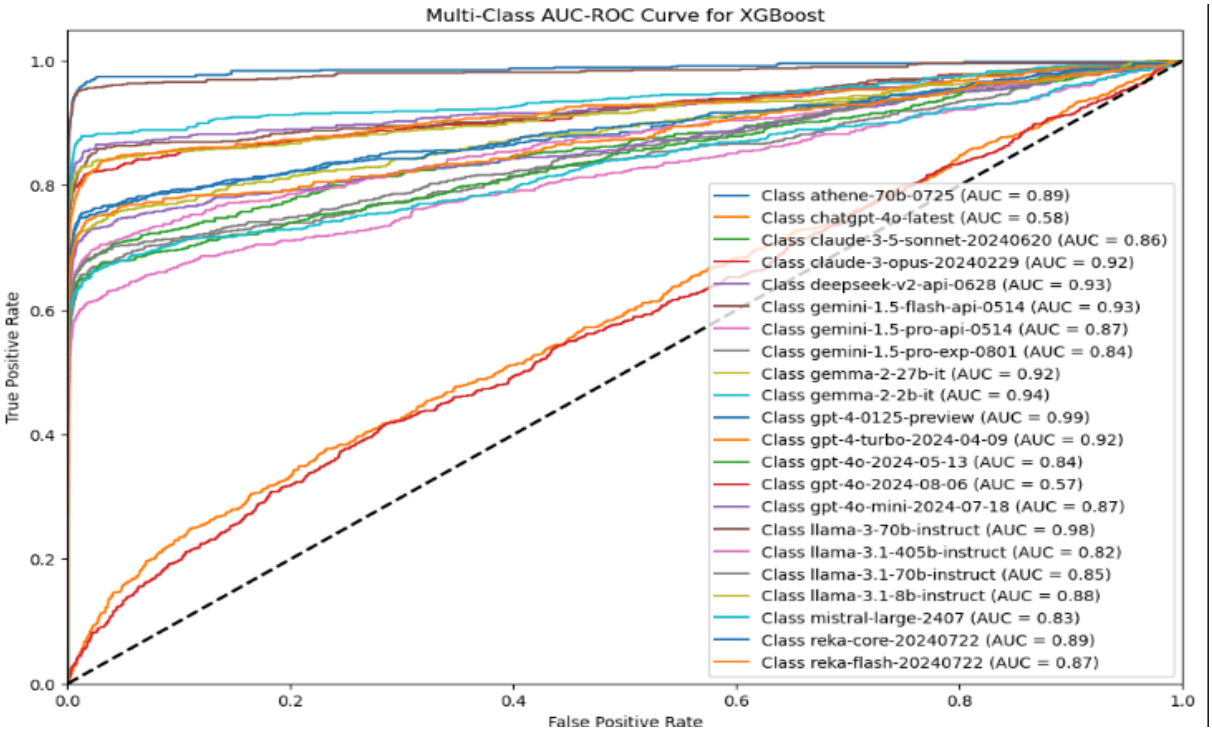

3. CATBOOST

The **CatBoost** classifier achieved an overall accuracy of 0.66.

Here we have used CatBoost and involve the following steps: Data Loading and Preprocessing, Hyperparameter Tuning using RandomizedSearchCV to tune hyperparameters and selecting the best parameters. Then it trains separate CatBoost classifiers One-vs-Rest approach. Making predictions which involves calculating overall accuracy, classification report for precision, recall, and F1-score, and visualizing the confusion matrix. Additionally, AUC-ROC curves are plotted.

```
Fitting 3 folds for each of 50 candidates, totalling 150 fits
Best Parameters for CatBoost: {'learning_rate': 0.1388888888888889, 'l2_leaf_reg': 3, 'iterations': 1000, 'depth': 10, 'border_count': 128, 'bagging_temperature': 3}
Overall Accuracy of Tuned CatBoost Classifier: 0.66
Classification Report:
                          precision    recall  f1-score   support

          athene-70b-0725       0.69      0.72      0.71       830
        chatgpt-4o-latest       0.16      0.10      0.12       790
claude-3-5-sonnet-20240620       0.62      0.64      0.63       776
    claude-3-opus-20240229       0.71      0.77      0.74       784
      deepseek-v2-api-0628       0.72      0.77      0.75       773
 gemini-1.5-flash-api-0514       0.74      0.78      0.76       821
   gemini-1.5-pro-api-0514       0.66      0.64      0.65       847
   gemini-1.5-pro-exp-0801       0.62      0.61      0.62       796
          gemma-2-27b-it       0.72      0.80      0.76       822
           gemma-2-2b-it       0.76      0.84      0.80       828
       gpt-4-0125-preview       0.79      0.92      0.85       747
    gpt-4-turbo-2024-04-09       0.67      0.76      0.72       799
       gpt-4o-2024-05-13       0.64      0.61      0.62       785
       gpt-4o-2024-08-06       0.13      0.08      0.10       769
    gpt-4o-mini-2024-07-18       0.68      0.65      0.67       840
      llama-3-70b-instruct       0.81      0.90      0.85       825
   llama-3.1-405b-instruct       0.62      0.58      0.60       807
    llama-3.1-70b-instruct       0.65      0.61      0.63       796
     llama-3.1-8b-instruct       0.71      0.69      0.70       797
...
                  accuracy                           0.66     17600
                 macro avg       0.64      0.66      0.65     17600
              weighted avg       0.64      0.66      0.65     17600
```

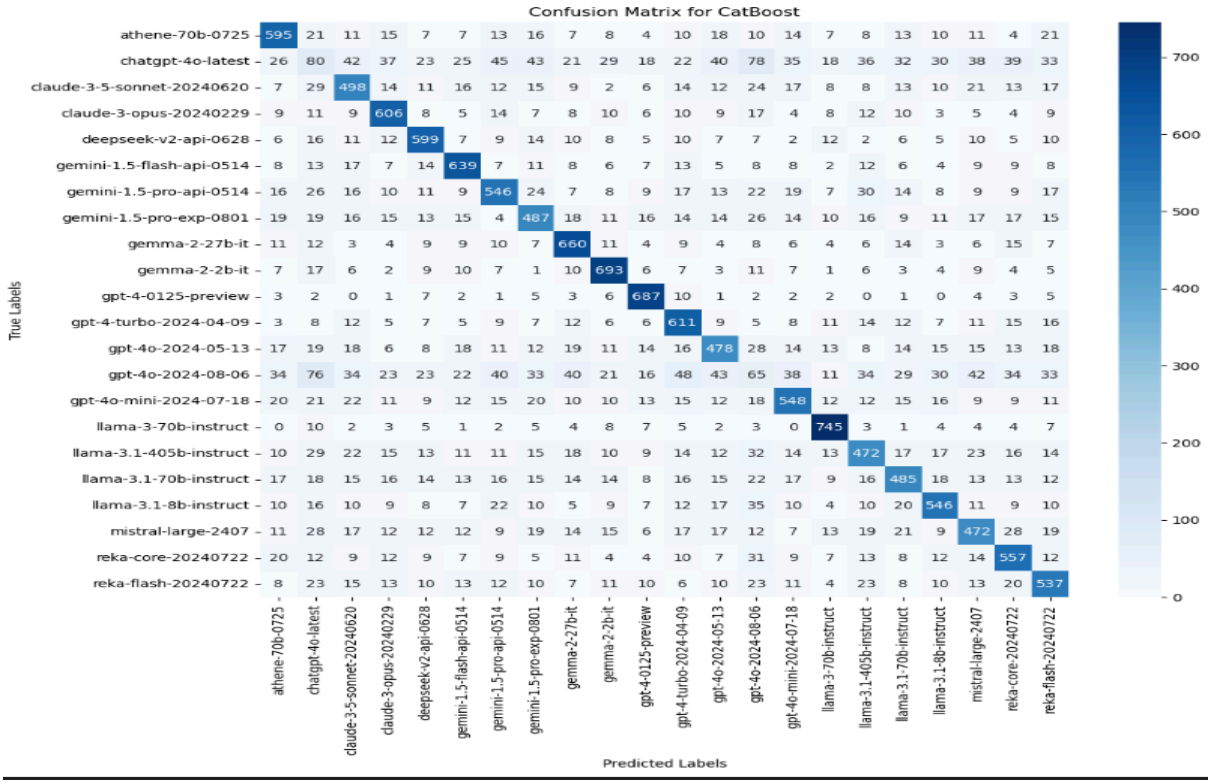

The Area Under the Curve (AUC) where values close to 1 indicate high ability and values closer to 0.5 suggest no better performance.

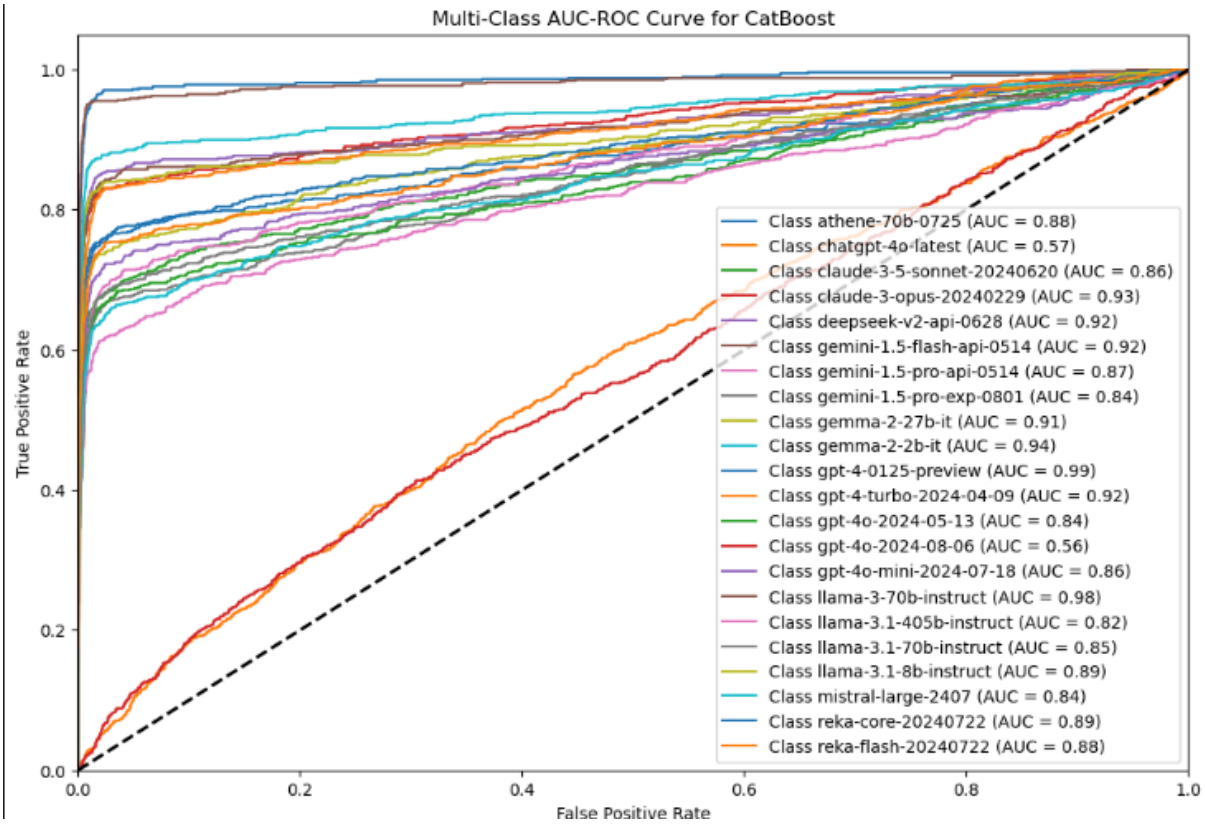

4. LIGHTGBM

The **LightGBM** classifier achieved an overall accuracy of 0.68.

The code uses LightGBM for a multi-class classification through a One-vs-Rest approach and involves the following steps: Data Loading and Preprocessing, Hyperparameter Tuning using RandomizedSearchCV with 150 iterations and 5-fold cross-validation to tune hyperparameters and selecting the best parameters. Then using the One-vs-Rest method, separate LightGBM classifiers and train it. Evaluating the results with confusion matrices, classification reports, and AUC-ROC curves.

```
Fitting 5 folds for each of 150 candidates, totalling 750 fits
Best Parameters for LightGBM: {'subsample': 0.5526315789473684, 'reg_lambda': 0, 'reg_alpha': 0, 'num_leaves': 150, 'n_estimators': 2000, 'min_split_gain'
Overall Accuracy of Tuned LightGBM Classifier: 0.68
Classification Report:
                         precision    recall  f1-score   support

          athene-70b-0725       0.73      0.73      0.73       830
        chatgpt-4o-latest       0.18      0.14      0.16       790
claude-3-5-sonnet-20240620      0.67      0.67      0.67       776
    claude-3-opus-20240229      0.72      0.80      0.76       784
      deepseek-v2-api-0628      0.78      0.82      0.80       773
 gemini-1.5-flash-api-0514      0.76      0.80      0.78       821
   gemini-1.5-pro-api-0514      0.67      0.67      0.67       847
   gemini-1.5-pro-exp-0801      0.64      0.62      0.63       796
           gemma-2-27b-it       0.78      0.81      0.80       822
            gemma-2-2b-it       0.82      0.85      0.83       828
       gpt-4-0125-preview       0.83      0.93      0.88       747
    gpt-4-turbo-2024-04-09      0.75      0.77      0.76       799
         gpt-4o-2024-05-13      0.66      0.64      0.65       785
         gpt-4o-2024-08-06      0.16      0.14      0.15       769
    gpt-4o-mini-2024-07-18      0.71      0.70      0.70       840
      llama-3-70b-instruct      0.85      0.93      0.89       825
   llama-3.1-405b-instruct      0.65      0.59      0.62       807
    llama-3.1-70b-instruct      0.67      0.64      0.65       796
     llama-3.1-8b-instruct      0.74      0.70      0.72       797
...
                  accuracy                           0.68     17600
                 macro avg      0.67      0.68      0.68     17600
              weighted avg      0.67      0.68      0.68     17600
```

High diagonal values suggest that the classifier performs well in predicting those classes, whereas high off-diagonal values indicate areas where the classifier confuses one class for another

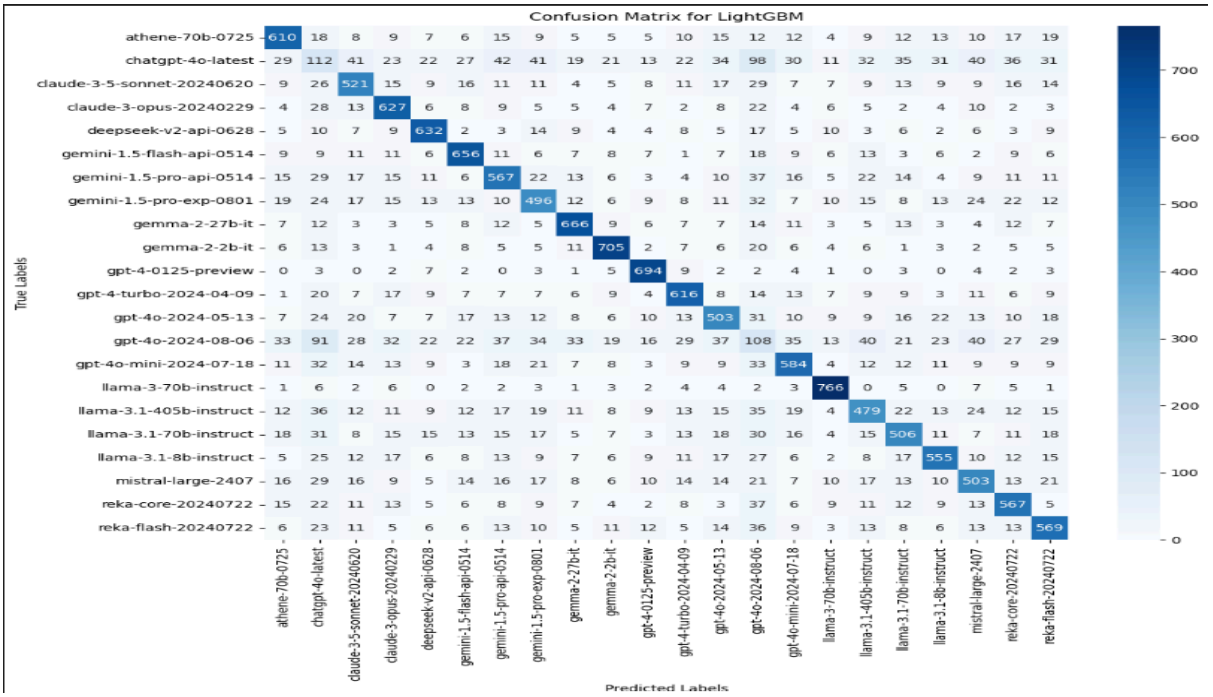

Values close to 1 indicating excellent ability and values near 0.5 suggesting no better accuracy.

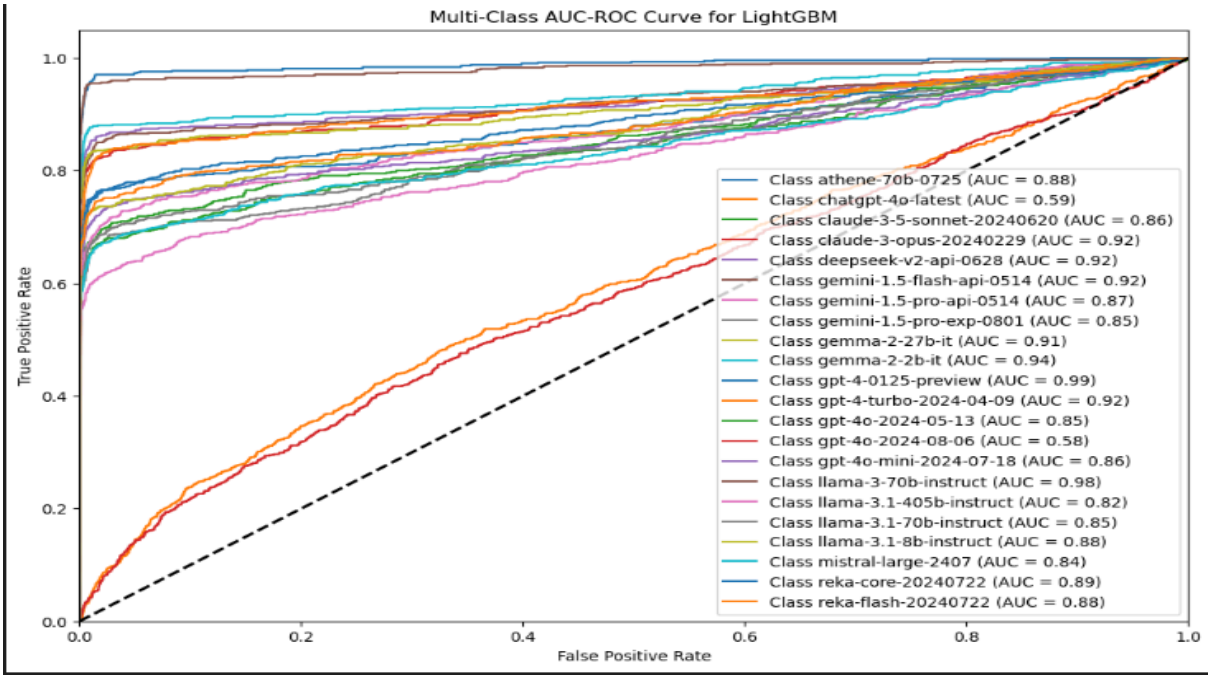

5. VOTING

The **Voting** classifier achieved an overall accuracy of 0.69.

We handled a multi-class classification by utilizing a **Voting Classifier** by combining LightGBM, CatBoost, and XGBoost through a One-vs-Rest (OvR) strategy. and involves the following steps: Data Loading and Preprocessing, Model Initialization: where it initializes three models LightGBM, CatBoost, and XGBoost with hyperparameters which constructs a Voting Classifier using soft voting from all three models. Then Train the One-vs-Rest Model with multiple Voting Classifiers in a One-vs-Rest. Evaluating the results with confusion matrices and classification reports.

```
Overall Accuracy of Voting Classifier: 0.69
Classification Report:
                           precision    recall  f1-score   support

          athene-70b-0725       0.71      0.73      0.72       830
        chatgpt-4o-latest       0.16      0.09      0.12       790
claude-3-5-sonnet-20240620      0.67      0.67      0.67       776
    claude-3-opus-20240229      0.71      0.80      0.75       784
      deepseek-v2-api-0628      0.76      0.82      0.79       773
 gemini-1.5-flash-api-0514      0.74      0.81      0.78       821
   gemini-1.5-pro-api-0514      0.69      0.67      0.68       847
   gemini-1.5-pro-exp-0801      0.66      0.64      0.65       796
           gemma-2-27b-it       0.77      0.82      0.79       822
            gemma-2-2b-it       0.80      0.85      0.83       828
       gpt-4-0125-preview       0.84      0.94      0.89       747
    gpt-4-turbo-2024-04-09      0.71      0.77      0.74       799
         gpt-4o-2024-05-13      0.65      0.65      0.65       785
         gpt-4o-2024-08-06      0.15      0.10      0.12       769
    gpt-4o-mini-2024-07-18      0.72      0.71      0.71       840
      llama-3-70b-instruct      0.82      0.93      0.87       825
   llama-3.1-405b-instruct      0.65      0.60      0.63       807
    llama-3.1-70b-instruct      0.65      0.65      0.65       796
     llama-3.1-8b-instruct      0.73      0.71      0.72       797
        mistral-large-2407      0.65      0.63      0.64       789
        reka-core-20240722      0.69      0.74      0.71       782
...
                  accuracy                           0.69     17600
                 macro avg      0.66      0.68      0.67     17600
              weighted avg      0.67      0.69      0.67     17600
```

The numbers on the matrix diagonal (e.g., 608 for the first class) indicate correct predictions for each class, while off-diagonal numbers represent misclassifications.

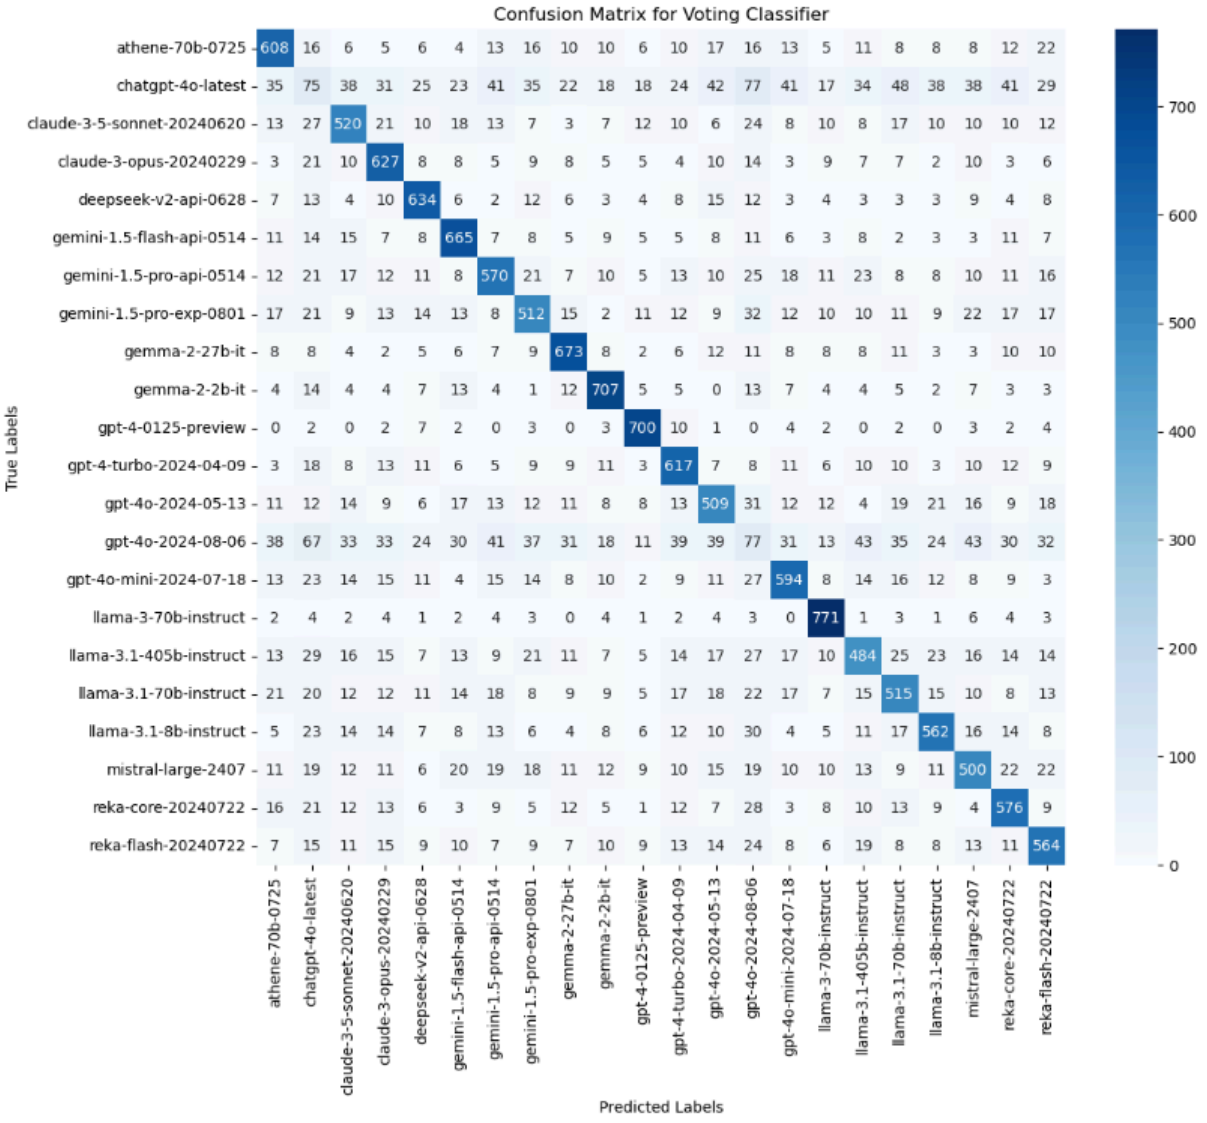

ACCURACY ACHIEVED BY EACH MODEL

MODEL NAME	ACCURACY
RANDOM FOREST	0.69
XG BOOST	0.68
CATBOOST	0.66
LIGHTGBM	0.68
VOTING	0.69